

Halo and more: exploring incremental verification and SNARKs without pairings

2021 Nov 05

[See all posts](#)

Special thanks to Justin Drake and Sean Bowe for wonderfully pedantic and thoughtful feedback and review, and to Pratyush Mishra for discussion that contributed to the original IPA exposition.

Readers who have been [following](#) the [ZK-SNARK](#) space [closely](#) should by now be familiar with the high level of how ZK-SNARKs work. ZK-SNARKs are based on checking equations where the elements going into the equations are mathematical abstractions like [polynomials](#) (or in [rank-1 constraint systems](#) matrices and vectors) that can hold a lot of data. There are three major families of cryptographic technologies that allow us to represent these abstractions succinctly: Merkle trees (for [FRI](#)), regular elliptic curves (for [inner product arguments \(IPAs\)](#)), and elliptic curves with pairings and trusted setups (for [KZG commitments](#)). These three technologies lead to the three types of proofs: FRI leads to STARKs, KZG commitments lead to "regular" SNARKs, and IPA-based schemes lead to bulletproofs. These three technologies have very distinct tradeoffs:

Technology	Cryptographic assumptions	Proof size	Verification time
FRI	Hashes only (quantum safe!)	Large (10-200 kB)	Medium (poly-logarithmic)
Inner product arguments (IPAs)	Basic elliptic curves	Medium (1-3 kB)	Very high (linear)
KZG commitments	Elliptic curves + pairings + trusted setup	Short (~500 bytes)	Low (constant)

So far, the first and the third have seen the most attention. The reason for this has to do with that pesky right column in the second row of the table: elliptic curve-based inner product arguments have linear verification time. What this means is that even though the *size* of a proof is small, the amount of time needed to verify the proof always takes longer than just running the computation yourself. This makes IPAs non-viable for scalability-related ZK-SNARK use cases: there's no point in using an IPA-based argument to prove the validity of an Ethereum block, because verifying the proof will take longer than just checking the block yourself. KZG and FRI-based proofs, on the other hand, really are much faster to verify than doing the computation yourself, so one of those two seems like the obvious choice.

More recently, however, there has been a slew of research into techniques for merging multiple IPA proofs into one. Much of the initial work on this was done as part of designing the [Halo protocol](#) which is [going into Zcash](#). These merging techniques are cheap, and a merged proof takes no longer to verify than a single one of the proofs that it's merging. This opens a way forward for IPAs to be useful: instead of verifying a size- (n) computation with a proof that takes still takes $(O(n))$ time to verify, break that computation up into smaller size- (k) steps, make $(\frac{n}{k})$ proofs for each step, and merge them together so the verifier's work goes down to a little more than $(O(k))$. These techniques also allow us to do *incremental verification*: if new things keep being introduced that need to be proven, you can just keep taking the existing proof, mixing it in with a proof of the new statement, and getting a proof of the new combined statement out. This is really useful for verifying the integrity of, say, an entire blockchain.

So how do these techniques work, and what can they do? That's exactly what this post is about.

Background: how do inner product arguments work?

Inner product arguments are a proof scheme that can work over many mathematical structures, but usually we focus on IPAs over [elliptic curve points](#). IPAs can be made over simple elliptic curves, theoretically even Bitcoin and Ethereum's [secp256k1](#) (though some special properties are preferred

to make [FFTs](#) more efficient); no need for insanely complicated pairing schemes that despite having written an [explainer article](#) and an [implementation](#) I can still barely understand myself.

We'll start off with the commitment scheme, typically called **Pedersen vector commitments**. To be able to commit to degree $(< n)$ polynomials, we first publicly choose a set of base points, $(G_0 \dots G_{n-1})$. These points can be generated through a pseudo-random procedure that can be re-executed by anyone (eg. the x coordinate of (G_i) can be $(\text{hash}(i, j))$ for the lowest integer $(j \geq 0)$ that produces a valid point); this is *not* a trusted setup as it does not rely on any specific party to introduce secret information.

To commit to a polynomial $(P(x) = \sum_i c_i x^i)$, the prover computes $(\text{com}(P) = \sum_i c_i G_i)$. For example, $(\text{com}(x^2 + 4))$ would equal $(G_2 + 4 * G_0)$ (remember, the $(+)$ and $(*)$ here are [elliptic curve addition](#) and multiplication). Cryptographers will also often add an extra $(r \cdot H)$ hiding parameter for privacy, but for simplicity of exposition we'll ignore privacy for now; in general, it's not that hard to add privacy into all of these schemes.



Though it's not really mathematically accurate to think of elliptic curve points as being like real numbers that have sizes, area is nevertheless a good intuition for thinking about linear combinations of elliptic curve points like we use in these commitments. The blue area here is the value of the Pedersen commitment $(C = \sum_i c_i G_i)$ to the polynomial $(P = \sum_i c_i x^i)$.

Now, let's get into how the proof works. **Our final goal will be a polynomial evaluation proof:** given some (z) , we want to make a proof that $(P(z) = a)$, where this proof can be verified by anyone who has the commitment $(C = \text{com}(P))$. **But first, we'll focus on a simpler task: proving that (C) is a valid commitment to any polynomial at all** - that is, proving that (C) was constructed by taking a linear combination $(\sum_i c_i G_i)$ of the points $(\{G_0 \dots G_{n-1}\})$, without anything else mixed in.

Of course, technically *any* point is some multiple of (G_0) and so it's theoretically a valid commitment of something, but what we care about is proving that the prover *knows* some $(\{c_0 \dots c_{n-1}\})$ such that $(\sum_i c_i G_i = C)$. A commitment (C) cannot commit to multiple distinct polynomials *that the prover knows about*, because if it could, that would imply that elliptic curves are broken.

The prover *could*, of course, just provide $(\{c_0 \dots c_{n-1}\})$ directly and let the verifier check the commitment. But this takes too much space. So instead, we try to reduce the problem to a smaller problem of half the size. The prover provides two points, (L) and (R) , representing the yellow and green areas in this diagram:



You may be able to see where this is going: if you add $(C + L + R)$ together (remember: (C) was the original commitment, so the blue area), the new combined point can be expressed as a sum of *four* squares instead of eight. And so now, the prover could finish by providing only four sums, the widths of each of the new squares. Repeat this protocol two more times, and we're down to a single full square, which the prover can prove by sending a single value representing its width.

But there's a problem: if (C) is incorrect in some way (eg. the prover added some extra point (H) into it), then the prover could just subtract (H) from (L) or (R) to compensate for it. We plug this hole by randomly scaling our points after the prover provides (L) and (R) :



Choose a random factor (α) (typically, we set (α) to be the hash of all data added to the proof so far, including the (L) and (R) , to ensure the verifier can also compute (α)). Every even (G_i) point gets scaled by (α) , every odd (G_i) point gets scaled *down* by the same factor. Every odd coefficient gets scaled up by (α) (notice the flip), and every even coefficient gets scaled down by (α) . Now, notice that:

- The yellow area (L) gets multiplied by (α^2) (because every yellow square is scaled up by (α) on both dimensions)
- The green area (R) gets divided by (α^2) (because every green square is scaled down by (α) on both dimensions)
- The blue area (C) remains unchanged (because its width is scaled up but its height is scaled down)

Hence, we can generate our new half-size instance of the problem with some simple transformations:

- $(G'_i) = \alpha G_{2i} + \frac{G_{2i+1}}{\alpha}$

- $(c'_i = \frac{c_{2i}}{\alpha} + \alpha c_{2i+1})$
- $(C' = C + \alpha^2 L + \frac{R}{\alpha^2})$

(Note: in some implementations you instead do $(G'_i = \alpha G_{2i} + G_{2i+1})$ and $(c'_i = c_{2i} + \alpha c_{2i+1})$ without dividing the odd points by (α) . This makes the equation $(C' = \alpha C + \alpha^2 L + R)$, which is less symmetric, but ensures that the function to compute any (G') in any round of the protocol becomes a polynomial without any division. [Yet another alternative](#) is to do $(G'_i = \alpha G_{2i} + G_{2i+1})$ and $(c'_i = c_{2i} + \frac{c_{2i+1}}{\alpha})$, which avoids any (α^2) terms.)

And then we repeat the process until we get down to one point:



Finally, we have a size-1 problem: prove that the final modified (C^*) (in this diagram it's (C''') because we had to do three iterations, but it's $(\log(n))$ iterations generally) equals the final modified (G^*_0) and (c^*_0) . Here, the prover just provides (c^*_0) in the clear, and the verifier checks $(c^*_0 G^*_0 = C^*)$. Computing (c^*_0) required being able to compute a linear combination of $(\{c_0 \dots c_{n-1}\})$ that was not known ahead of time, so providing it and verifying it convinces the verifier that the prover actually does know all the coefficients that go into the commitment. This concludes the proof.

Recapping:

- The statement we are proving is that (C) is a commitment to *some* polynomial $(P(x) = \sum_i c_i x^i)$ committed to using the agreed-upon base points $(\{G_0 \dots G_{n-1}\})$
- The proof consists of $(\log(n))$ pairs of (L, R) values, representing the yellow and green areas at each step. The prover also provides the final (c^*_0)
- The verifier walks through the proof, generating the (α) value at each step using the same algorithm as the prover and computing the new (C') and (G'_i) values (the verifier doesn't know the (c_i) values so they can't compute any (c'_i) values)
- At the end, they check whether or not $(c^*_0 G^*_0 = C^*)$


On the whole, the proof contains $(2 * \log(n))$ elliptic curve points and one number (for pedants: one *field element*). Verifying the proof takes logarithmic time in every step except one: computing the new (G'_i) values. This step is, unfortunately, linear.

See also: [Dankrad Feist's more detailed explanation](#) of inner product arguments.

Extension to polynomial evaluations

We can extend to polynomial evaluations with a simple clever trick. Suppose we are trying to prove $(P(z) = a)$. The prover and the verifier can extend the base points $(G_0 \dots G_{n-1})$ by attaching powers of (z) to them: the new base points become $((G_0, 1), (G_1, z) \dots (G_{n-1}, z^{n-1}))$. These pairs can be treated as mathematical objects (for pedants: *group elements*) much like elliptic curve points themselves; to add them you do so element-by-element: $((A, x) + (B, y) = ((A + B, x + y))$, using elliptic curve addition for the points and regular field addition for the numbers.

We can make a Pedersen commitment using this extended base!



Now, here's a puzzle. Suppose $(P(x) = \sum_i c_i x^i)$, where $(P(z) = a)$, would have a commitment $(C = \sum_i c_i G_i)$ if we were to use the regular elliptic curve points we used before as a base. If we use the pairs $((G_i, z^i))$ as a base instead, the commitment would be $((C, y))$ for some (y) . What must be the value of (y) ?

The answer is: it must be equal to (a) ! This is easy to see: the commitment is $((C, y) = (\sum_i c_i (G_i, z^i)))$, which we can decompose as $((\sum_i c_i G_i, \sum_i c_i z^i))$. The former is equal to (C) , and the latter is just the evaluation $(P(z))$!

Hence, if (C) is a "regular" commitment to (P) using $(G_0 \dots G_{n-1})$ as a base, then to prove that $(P(z) = a)$ we need only use the same protocol above, but proving that $((C, a))$ is a valid commitment using $((G_0, 1), (G_1, z) \dots (G_{n-1}, z^{n-1}))$ as a base!

Note that in practice, this is usually done slightly differently as an optimization: instead of attaching the numbers to the points and explicitly dealing with structures of the form $((G_i, z^i))$, we add another randomly chosen base point (H) and express it as $(G_i + z^i H)$. This saves space.

See [here](#) for an example implementation of this whole protocol.

So, how do we combine these proofs?

Suppose that you are given two polynomial evaluation proofs, with different polynomials and different evaluation points, and want to make a proof that they are both correct. You have:

- Proof (Π_1) proving that $(P_1(z_1) = y_1)$, where (P_1) is represented by $(\text{com}(P_1) = C_1)$
- Proof (Π_2) proving that $(P_2(z_2) = y_2)$, where (P_2) is represented by $(\text{com}(P_2) = C_2)$

Verifying each proof takes linear time. We want to make a proof that proves that both proofs are correct. This will still take linear time, but the verifier will only have to make *one* round of linear time verification instead of two.

We start off with an observation. The only linear-time step in performing the verification of the proofs is computing the (G'_i) values. This is $(O(n))$ work because you have to combine $(\frac{n}{2})$ pairs of (G_i) values into (G'_i) values, then $(\frac{n}{4})$ pairs of (G'_i) values into (G''_i) values, and so on, for a total of (n) combinations of pairs. But if you look at the algorithm carefully, you will notice that we *don't actually need any of the intermediate (G'_i) values; we only need the final (G^{*}_0)* . This (G^{*}_0) is a linear combination of the initial (G_i) values. What are the

coefficients to that linear combination? It turns out that the (G_i) coefficient is the (X^i) term of this polynomial:

$$[(X + \alpha_1) * (X^2 + \alpha_2) * \dots * (X^{\frac{n}{2}} + \alpha_{\log(n)})]$$

This is using the $(C' = \alpha C + \alpha^2 L + R)$ version we mentioned above. The ability to directly compute (G^*_0) as a linear combination already cuts down our work to $(O(\frac{n}{\log(n)}))$ due to [fast linear combination algorithms](#), but we can go further.

The above polynomial has degree $(n - 1)$, with (n) nonzero coefficients. But its un-expanded form has size $(\log(n))$, and so you can *evaluate* the polynomial at any point in $(O(\log(n)))$ time. Additionally, you might notice that (G^*_0) is a commitment to this polynomial, so we can directly *prove* evaluations! So here is what we do:

- The prover computes the above polynomial for each proof; we'll call these polynomials (K_1) with $(\text{com}(K_1) = D_1)$ and (K_2) with $(\text{com}(K_2) = D_2)$. In a "normal" verification, the verifier would be computing (D_1) and (D_2) themselves as these are just the (G^*_0) values for their respective proofs. Here, the prover provides (D_1) and (D_2) and the rest of the work is proving that they're correct.
- To prove the correctness of (D_1) and (D_2) we'll prove that they're correct at a random point. We choose a random point (t) , and evaluate both $(e_1 = K_1(t))$ and $(e_2 = K_2(t))$
- The prover generates a random linear combination $(L(x) = K_1(x) + rK_2(x))$ (and the verifier can generate $(\text{com}(L) = D_1 + rD_2)$). The prover now just needs to make a single proof that $(L(t) = e_1 + re_2)$.



The verifier still needs to do a bunch of extra steps, but all of those steps take either $(O(1))$ or $(O(\log(n)))$ work: evaluate $(e_1 = K_1(t))$ and $(e_2 = K_2(t))$, calculate the (α_i) coefficients of both (K_i) polynomials in the first place, do the elliptic curve addition $(\text{com}(L) = D_1 + rD_2)$. But this all takes vastly less than linear time, so all in all we still benefit: the verifier only needs to do the linear-time step of computing a (G^*_0) point themselves once.

This technique can easily be generalized to merge $(m > 2)$ signatures.

From merging IPAs to merging IPA-based SNARKs: Halo

Now, we get into the core mechanic of the [Halo protocol](#) being integrated in Zcash, which uses this proof combining technique to create a recursive proof system. The setup is simple: suppose you have a chain, where each block has an associated IPA-based SNARK (see [here](#) for how generic SNARKs from polynomial commitments work) proving its correctness. You want to create a new block, building on top of the previous tip of the chain. The new block should have its own IPA-based SNARK proving the correctness of the block. In fact, this proof should cover both the correctness of the new block *and* the correctness of the previous block's proof of the correctness of the entire chain before it.

IPA-based proofs by themselves cannot do this, because a proof of a statement takes longer to verify than checking the statement itself, so a proof of a proof will take even longer to verify than both proofs separately. But proof merging can do it!



Essentially, we use the usual "recursive SNARK" technique to verify the proofs, except the "proof of a proof" part is only proving the logarithmic part of the work. We add an extra chain of aggregate proofs, using a trick similar to the proof merging scheme above, to handle the linear part of the new work. To verify the whole chain, the verifier need only verify one linear-time proof at the very tip of the chain.

The precise details are somewhat different from the exact proof-combining trick in the previous section for efficiency reasons. Instead of using the proof-combining trick to combine multiple proofs, we use it on a *single* proof, just to re-randomize the point that the polynomial committed to by (G^*_0) needs to be evaluated at. We then use the *same* newly chosen evaluation point to evaluate the polynomials in the proof of the block's correctness, which allows us to prove the polynomial evaluations together in a single IPA.

Expressed in math:

- Let $(P(z) = a)$ be the previous statement that needs to be proven
- The prover generates (G^*_0)
- The prover proves the correctness of the new block plus the logarithmic work in the previous statements by generating a [PLONK proof](#): $(Q_L * A + Q_R * B + Q_O * C + Q_M * A * B + Q_C = Z * H)$
- The prover chooses a random point (t) , and proves the evaluation of a linear combination of $(\{G^*_0, Q_L, A, Q_R, B, Q_O, C, Q_M, Q_C, Z, H\})$ at (t) . We can then check the above equation, replacing each polynomial with its now-verified evaluation at (t) , to verify the PLONK proof.

Incremental verification, more generally

The size of each "step" does not need to be a full block verification; it could be something as small as a single step of a virtual machine. The smaller the steps the better: it ensures that the linear work that the verifier ultimately has to do at the end is less. The only lower bound is that each step has to be big enough to contain a SNARK verifying the $(\log(n))$ portion of the work of a step.

But regardless of the fine details, this mechanism allows us to make succinct and easy-to-verify SNARKs, including easy support for recursive proofs that allow you to extend proofs in real time as the computation extends and even have different provers to do different parts of the proving work,

all without pairings or a trusted setup! The main downside is some extra technical complexity, compared with a "simple" polynomial-based proof using eg. KZG-based commitments.

Technology	Cryptographic assumptions	Proof size	Verification time
FRI	Hashes only (quantum safe!)	Large (10-200 kB)	Medium (poly-logarithmic)
Inner product arguments (IPAs)	Basic elliptic curves	Medium (1-3 kB)	Very high (linear)
KZG commitments	Elliptic curves + pairings + trusted setup	Short (~500 bytes)	Low (constant)
IPA + Halo-style aggregation	Basic elliptic curves	Medium (1-3 kB)	Medium (constant but higher than KZG)

Not just polynomials! Merging R1CS proofs

A common alternative to building SNARKs out of polynomial games is building SNARKs out of matrix-vector multiplication games. Polynomials and vectors+matrices are both natural bases for SNARK protocols because they are mathematical abstractions that can store and compute over large amounts of data at the same time, and that admit commitment schemes that allow verifiers to check equations quickly.

In R1CS (see a more detailed description [here](#)), an instance of the game consists of three matrices A , B , C , and a solution is a vector Z such that $(A \cdot Z) \circ (B \cdot Z) = C \cdot Z$ (the problem is often in practice restricted further by requiring the prover to make part of Z public and requiring the last entry of Z to be 1).

A		B		C	
1	5	1	1	1	0
3	0	3	0	3	0
35	0	35	0	35	1
9	0	9	0	9	0
27	0	27	0	27	0
30	1	30	0	30	0

$35 \quad * \quad 1 \quad - \quad 35 \quad = \quad 0$

An R1CS instance with a single constraint (so A , B and C have width 1), with a satisfying Z vector, though notice that here the Z appears on the left and has 1 in the top position instead of the bottom.

Just like with polynomial-based SNARKs, this R1CS game can be turned into a proof scheme by creating commitments to A , B and C , requiring the prover to provide a commitment to (the private portion of) Z , and using fancy proving tricks to prove the equation $(A \cdot Z) \circ (B \cdot Z) = C \cdot Z$, where \circ is item-by-item multiplication, without fully revealing any of these objects. And just like with IPAs, this R1CS game has a proof merging scheme!

Ioanna Tzialla et al describe such a scheme in a [recent paper](#) (see page 8-9 for their description). They first modify the game by introducing an expanded equation:

$$[(A \cdot Z) \circ (B \cdot Z) - u \cdot (C \cdot Z) = E]$$

For a "base" instance, $(u = 1)$ and $(E = 0)$, so we get back the original R1CS equation. The extra slack variables are added to make aggregation possible; aggregated instances will have other values of (u) and (E) . Now, suppose that you have two solutions to the same instance, though with different (u) and (E) variables:

$$(A \cdot Z_1) \cdot (B \cdot Z_1) - u_1 \cdot (C \cdot Z_1) = E_1$$

$$(A \cdot Z_2) \cdot (B \cdot Z_2) - u_2 \cdot (C \cdot Z_2) = E_2$$

The trick involves taking a random linear combination $(Z_3 = Z_1 + r Z_2)$, and making the equation work with this new value. First, let's evaluate the left side:

$$(A \cdot (Z_1 + rZ_2)) \cdot (B \cdot (Z_1 + rZ_2)) - (u_1 + ru_2) \cdot (C \cdot (Z_1 + rZ_2))$$

This expands into the following (grouping the (1) , (r) and (r^2) terms together):

$$(A \cdot Z_1) \cdot (B \cdot Z_1) - u_1 \cdot (C \cdot Z_1)$$

$$r(A \cdot Z_1) \cdot (B \cdot Z_2) + (A \cdot Z_2) \cdot (B \cdot Z_1) - u_1 \cdot (C \cdot Z_2) - u_2 \cdot (C \cdot Z_1)$$

$$r^2(A \cdot Z_2) \cdot (B \cdot Z_2) - u_2 \cdot (C \cdot Z_2)$$

The first term is just (E_1) ; the third term is $(r^2 \cdot E_2)$. The middle term is very similar to the cross-term (the yellow + green areas) near the very start of this post. The prover simply provides the middle term (without the (r) factor), and just like in the IPA proof, the randomization forces the prover to be honest.

Hence, it's possible to make merging schemes for R1CS-based protocols too. Interestingly enough, we don't even technically need to have a "succinct" protocol for proving the $(A \cdot Z) \cdot (B \cdot Z) = u \cdot (C \cdot Z) + E$ relation at the end; instead, the prover could just prove by opening all the commitments directly! This would still be "succinct" because the verifier would only need to verify one proof that actually represents an arbitrarily large number of statements. However, in practice having a succinct protocol for this last step is better because it keeps the proofs smaller, and [Tzialla et al's paper](#) provides such a protocol too (see page 10).

Recap

- We don't know of a way to make a commitment to a size- (n) polynomial where evaluations of the polynomial can be verified in $(< O(n))$ time directly. The best that we can do is make a $(\log(n))$ sized proof, where all of the work to verify it is logarithmic *except for one final $(O(n))$ -time piece*.
- But what we *can* do is merge multiple proofs together. Given (m) proofs of evaluations of size- (n) polynomials, you can make a proof that covers *all* of these evaluations, that takes logarithmic work plus a single size- (n) polynomial proof to verify.
- With some clever trickery, separating out the logarithmic parts from the linear parts of proof verification, we can leverage this to make recursive SNARKs.
- These recursive SNARKs are actually more efficient than doing recursive SNARKs "directly"! In fact, even in contexts where direct recursive SNARKs are possible (eg. proofs with KZG commitments), Halo-style techniques are typically used instead because they are more efficient.
- It's not just about polynomials; other games used in SNARKs like R1CS can also be aggregated in similar clever ways.
- No pairings or trusted setups required!

The march toward faster and more efficient and safer ZK-SNARKs just keeps going...