

Gitcoin Grants Round 5 Retrospective

2020 Apr 30

[See all posts](#)

Special thanks to Kevin Owocki and Frank Chen for help and review

Round 5 of Gitcoin Grants has just finished, with \$250,000 of matching split between tech, media, and the new (non-Ethereum-centric) category of "public health". In general, it seems like the mechanism and the community are settling down into a regular rhythm. People know what it means to contribute, people know what to expect, and the results emerge in a relatively predictable pattern - even if which specific grants get the most funds is not so easy to predict.

CLR MATCHING ROUND 5			
Tech Grants			
	NUMBER OF CONTRIBUTIONS	TOTAL CONTRIBUTED	CLR MATCHING
1 White Hat Hacking	234	\$4,687	\$15,704
2 Arboreum	178	\$7,038	\$9,046
3 1inch.exchange	180	\$50,712	\$7,893
4 The Commons Stack	145	\$3,061	\$6,497
5 POAP	128	\$2,940	\$5,139
6 Tornado.cash	142	\$1,178	\$3,496
7 DAppNode	107	\$1,993	\$3,352
8 Rotki	111	\$1,578	\$3,165
9 Gitcoin Grants Dev	116	\$2,688	\$3,087
10 Prysm by Prysmatic Labs	119	\$1,939	\$3,055

CLR MATCHING ROUND 5			
Media Grants			
	NUMBER OF CONTRIBUTIONS	TOTAL CONTRIBUTED	CLR MATCHING
1 Week in Ethereum News	157	\$2,141	\$10,054
2 Chris Blec	143	\$1,390	\$5,716
3 EthHub	123	\$1,566	\$5,148
4 The Defiant	116	\$3,748	\$4,886
5 MetaCartel	99	\$610	\$3,232
6 Transaction Pending	41	\$4,358	\$2,547
7 Bankless (Translations)	81	\$1,096	\$2,167
8 @antiprosynth	68	\$1,576	\$1,897
9 DeFi Dad Tutorials	73	\$465	\$1,793
10 Interspace.chat	54	\$1,007	\$1,755

Stability of income

So let's go straight into the analysis. One important property worth looking at is stability of income across rounds: do projects that do well in round N also tend to do well in round N+1? Stability of income is very important if we want to support an ecosystem of "quadratic freelancers": we want people to feel comfortable relying on their income knowing that it will not completely disappear the next round. On the other hand, it would be harmful if some recipients became completely entrenched, with no opportunity for new projects to come in and compete for the pot, so there is a need for a balance.

On the media side, we do see some balance between stability and dynamism:

Round 4 vs Round 5 income (media)



Week in Ethereum had the highest total amount received in both [the previous round](#) and the current round. EthHub and Bankless are also near the top in both the current round and the previous round. On the other hand, Antiprosynthesis, the (beloved? notorious? famous?) Twitter info-warrior, has decreased from \$13,813 to \$5,350, while [Chris Blec's YouTube channel](#) has *increased* from \$5,851 to \$12,803. So some churn, but also some continuity between rounds.

On the tech side, we see much more churn in the winners, with a less clear relationship between income last round and income this round:

Round 4 vs Round 5 income (tech)



Last round, the winner was Tornado Cash, claiming \$30,783; this round, they are down to \$8,154. This round, the three roughly-even winners are [Samczsun](#) (\$4,631 contributions + \$15,704 match = \$20,335 total), [Arboreum](#) (\$16,084 contributions + \$9,046 match = \$25,128 total) and [linch.exchange](#) (\$58,566 contributions + \$7,893 match = \$66,459 total), in the latter case the bulk coming from one contribution:

06 Apr 2020


tgerring
47,500.0000 DAI
(+2500.0 DAI optional tip to [Gitcoin](#))

47,500.0000 DAI

In the previous round, those three winners were not even in the top ten, and in some cases not even part of Gitcoin Grants at all.

These numbers show us two things. First, large parts of the Gitcoin community seem to be in the mindset of treating grants not as a question of "how much do you deserve for your last two months of work?", but rather as a one-off reward for years of contributions in the past. This was one of the strongest rebuttals that I received to my [criticism of Antiprosynthesis receiving \\$13,813 in the last round](#): that the people who

contributed to that award did not see it as two months' salary, but rather as a reward for years of dedication and work for the Ethereum ecosystem. In the next round, contributors were content that the debt was sufficiently repaid, and so they moved on to give a similar gift of appreciation and gratitude to Chris Blec.

That said, not everyone contributes in this way. For example, Prysm got \$7,966 last round and \$8,033 this round, and Week in Ethereum is consistently well-rewarded (\$16,727 previous, \$12,195 current), and EthHub saw less stability but still kept half its income (\$13,515 previous, \$6,705 current) even amid a 20% drop to the matching pool size as some funds were redirected to public health. So there definitely are some contributors that *are* getting almost a reasonable monthly salary from Gitcoin Grants (yes, even these amounts are all serious underpayment, but remember that the pool of funds Gitcoin Grants has to distribute in the first place is quite small, so there's no allocation that would *not* seriously underpay most people; the hope is that in the future we will find ways to make the matching pot grow bigger).

Why didn't more people use recurring contributions?

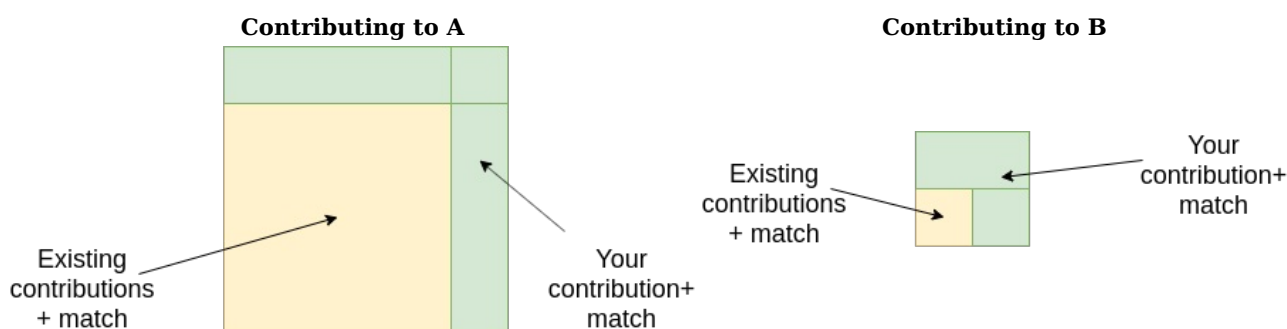
One feature that was tested this round to try to improve stability was recurring contributions: users could choose to split their contribution among multiple rounds. However, the feature was not used often: out of over 8,000 total contributions, only 120 actually made recurring contributions. I can think of three possible explanations for this:

1. People just don't want to give recurring contributions; they genuinely prefer to freshly rethink who they are supporting every round.
2. People would be willing to give more recurring contributions, but there is some kind of "market failure" stopping them; that is, it's collectively optimal for everyone to give more recurring contributions, but it's not any individual contributor's interest to be the first to do so.
3. There's some UI inconveniences or other "incidental" obstacles preventing recurring contributions.

In a recent call with the Gitcoin team, hypothesis (3) was mentioned frequently. A specific issue was that people were worried about making recurring contributions because they were concerned whether or not the money that they lock up for a recurring contribution would be safe. Improving the payment system and notification workflow may help with this. Another option is to move away from explicit "streaming" and instead simply have the UI provide an option for repeating the last round's contributions and making edits from there.

Hypothesis (1) also should be taken seriously; there's genuine value in preventing ossification and allowing space for new entrants. But I want to zoom in particularly on hypothesis (2), the coordination failure hypothesis.

My explanation of hypothesis (2) starts, interestingly enough, with a defense of (1): why ossification is genuinely a risk. Suppose that there are two projects, A and B, and suppose that they are equal quality. But A already has an established base of contributors; B does not (we'll say for illustration it only has a few existing contributors). Here's how much matching you are contributing by participating in each project:



Clearly, you have more impact by supporting A, and so A gets even more contributors and B gets fewer; the rich get richer. Even if project B was *somewhat better*, the greater impact from supporting A could still create a lock-in that reinforces A's position. The current everyone-starts-from-zero-in-each-round mechanism greatly limits this type of entrenchment, because, well, everyone's matching gets reset and starts from zero.

However, a very similar effect also is the cause behind the market failure preventing stable recurring contributions, and the every-round-reset *actually exacerbates it*. Look at the same picture above, except instead of thinking of A and B as *two different projects*, think of them as *the same project in the current round and in the next round*.

We simplify the model as follows. An individual has two choices: contribute \$10 in the current round, or contribute \$5 in the current round and \$5 in the next round. If the matchings in the two rounds were equal, then the latter option would actually be more favorable: because the matching is proportional to the square root of the donation size, the former might give you eg. a \$200 match now, but the latter would give you \$141 in the current round + \$141 in the next round = \$282. But if you see a large mass of people contributing in the current round, and you expect much fewer people to contribute in the second round, then the choice is not \$200 versus \$141 + \$141, it might be \$200 versus \$141 + \$5. And so you're better off joining the current

round's frenzy. We can mathematically analyze the equilibrium:



So there is a substantial region within which the bad equilibrium of everyone concentrating is sticky: if more than about 3/4 of contributors are expected to concentrate, it seems in your interest to also concentrate. A mathematically astute reader may note that there is [always some intermediate strategy](#) that involves splitting but at a ratio different from 50/50, which you can prove performs better than either full concentrating *or* the even split, but here we get back to hypothesis (3) above: the UI doesn't offer such a complex menu of choices, it just offers the choice of a one-time contribution or a recurring contribution, so people pick one or the other.

How might we fix this? One option is to add a bit of continuity to matching ratios: when computing pairwise matches, match against not just the current round's contributors but, say, 1/3 of the previous round's contributors as well:



This makes some philosophical sense: the objective of quadratic funding is to subsidize contributions to projects that are detected to be public goods because multiple people have contributed to them, and contributions in the previous round are certainly also evidence of a project's value, so why not reuse those? So here, moving away from everyone-starts-from-zero toward this partial carryover of matching ratios would mitigate the round concentration effect - but, of course, it would exacerbate the risk of entrenchment. Hence, some experimentation and balance may be in order. A broader philosophical question is, is there really a deep inherent tradeoff between risk of entrenchment and stability of income, or is there some way we could get both?

Responses to negative contributions

This round also introduced negative contributions, a feature proposed in my [review of the previous round](#). But as with recurring contributions, very few people made negative contributions, to the point where their impact on the results was negligible. Also, there was [active opposition](#) to [negative contributions](#):



Source: honestly I have no idea, someone else sent it to me and they forgot where they found it. Sorry :(

The main source of opposition was basically what I predicted in the previous round. Adding a mechanism that allows people to penalize others, even if deservedly so, can have tricky and easily harmful social consequences. Some people even opposed the negative contribution mechanism to the point where they took care to give positive contributions to everyone who received a negative contribution.

How do we respond? To me it seems clear that, in the long run, *some* mechanism of filtering out bad projects, and ideally compensating for overexcitement into good projects, will have to exist. It doesn't necessarily need to be integrated as a symmetric part of the QF, but there does need to be a filter of some form. And this mechanism, whatever form it will take, invariably opens up the possibility of the same social dynamics. So there is a challenge that will have to be solved no matter how we do it.

One approach would be to hide more information: instead of just hiding *who* made a negative contribution, outright hide the fact that a negative contribution was made. Many opponents of negative contributions explicitly indicated that they would be okay (or at least more okay) with such a model. And indeed (see the next section), this is a direction we will have to go anyway. But it would come at a cost - effectively hiding negative contributions would mean not giving as much real-time feedback into what projects got how much funds.

Stepping up the fight against collusion

This round saw much larger-scale attempts at collusion:




owoki, GItter of Coins
@owocki



Today I have a major decision to announce.

TLDR - We have identified an instance of collusion in the health round and are counting contributions we identify as colluding as being from the same account.

Next time we will have a more robust identity system.

Details 

12:23 PM · Apr 14, 2020 · [Twitter Web App](#)

15 Retweets 78 Likes

It does seem clear that, at current scales, stronger protections against manipulation are going to be required. The first thing that can be done is adding a stronger identity verification layer than Github accounts; this is something that the Gitcoin team is already working on. There is definitely a complex tradeoff between security and inclusiveness to be worked through, but it is not especially difficult to implement a first version. And if the identity problem is solved to a reasonable extent, that will likely be enough to prevent collusion at current scales. But in the longer term, we are going to need protection not just against manipulating the system by making many fake accounts, but also against collusion via bribes (explicit and implicit).



[MACI](#) is the solution that I proposed (and Barry Whitehat and co are implementing) to solve this problem. Essentially, MACI is a cryptographic construction that allows for contributions to projects to happen on-chain in a privacy-preserving, encrypted form, that allows anyone to cryptographically verify that the mechanism is being implemented *correctly*, but prevents participants from being able to prove to a third party that they made any particular contribution. Unprovability means that if someone tries to bribe others to contribute to their project, the bribe recipients would have no way to prove that they actually contributed to that project, making the bribe unenforceable. Benign "collusion" in the form of friends and family supporting each other would still happen, as people would not easily lie to each other at such small scales, but any broader collusion would be very difficult to maintain.

However, we do need to think through some of the second-order consequences that integrating MACI would introduce. The biggest blessing, and curse, of using MACI is that contributions become hidden. Identities necessarily become hidden, but even the exact timing of contributions would need to be hidden to prevent deanonymization through timing (to prove that *you* contributed, make the total amount jump up between 17:40 and 17:42 today). Instead, for example, totals could be provided and updated once per day. Note that as a corollary negative contributions would be hidden as well; they would only appear if they exceeded all positive contributions for an entire day (and if even that is not desired then the mechanism for when balances are updated could be tweaked to further hide downward changes).

The challenge with hiding contributions is that we lose the "social proof" motivator for contributing: if contributions are unprovable you can't as easily publicly brag about a contribution you made. My best proposal for solving this is for the mechanism to publish one extra number: the *total* amount that a particular participant contributed (counting only projects that have received at least 10 contributors to prevent inflating one's number by self-dealing). Individuals would then have a generic "proof-of-generosity" that they contributed some specific *total* amount, and could publicly state (without proof) what projects it was that they supported. But this is all a significant change to the user experience that will require multiple rounds of experimentation to get right.

Conclusions

All in all, Gitcoin Grants is establishing itself as a significant pillar of the Ethereum ecosystem that more and more projects are relying on for some or all of their support. While it has a relatively low amount of funding at present, and so inevitably underfunds almost everything it touches, we hope that over time we'll continue to see larger sources of funding for the matching pools appear. One option is [MEV auctions](#), another is that new or existing token projects looking to do airdrops could provide the tokens to a matching pool. A third is transaction fees of various applications. With larger amounts of funding, Gitcoin Grants could serve as a more significant funding stream - though to get to that point, further iteration and work on fine-tuning the mechanism will be required.

CLR MATCHING ROUND 5			
 Health Grants			
		NUMBER OF CONTRIBUTIONS	TOTAL CONTRIBUTED
			CLR MATCHING
1	CIC (COVID-19) Kenyan Crisis Aid	85	\$2,767
2	Open Source Covid Ventilators + Masks	89	\$2,472
3	Mask + Test Kit Mutual Aid Fund	87	\$2,447
4	CuraDAO COVID-19 Campaign	75	\$1,056
5	Giveth & Coz: Giving to COVID-19 Causes	57	\$6,027
6	African Angels	48	\$1,970
7	Decentralised supply chain of 3D printed protective gear	49	\$700
8	Tracy	93	\$3,186
9	Save the Children	52	\$6,963
10	Collab19	35	\$175

Additionally, this round saw Gitcoin Grants' first foray into applications beyond Ethereum with the health section. There is growing interest in quadratic funding from local government bodies and other non-blockchain groups, and it would be very valuable to see quadratic funding more broadly deployed in such contexts. That said, there are unique challenges there too. First, there's issues around onboarding people who do not already have cryptocurrency. Second, the Ethereum community is naturally expert in the needs of the Ethereum community, but neither it nor average people are expert in, eg. medical support for the coronavirus pandemic. We should expect quadratic funding to perform worse when the participants are not experts in the domain they're being asked to contribute to. Will non-blockchain uses of QF focus on domains where there's a clear local community that's expert in its own needs, or will people try larger-scale deployments soon? If we do see larger-scale deployments, how will those turn out? There's still a lot of questions to be answered.